

## Concurrency & IPC in Kernel:-

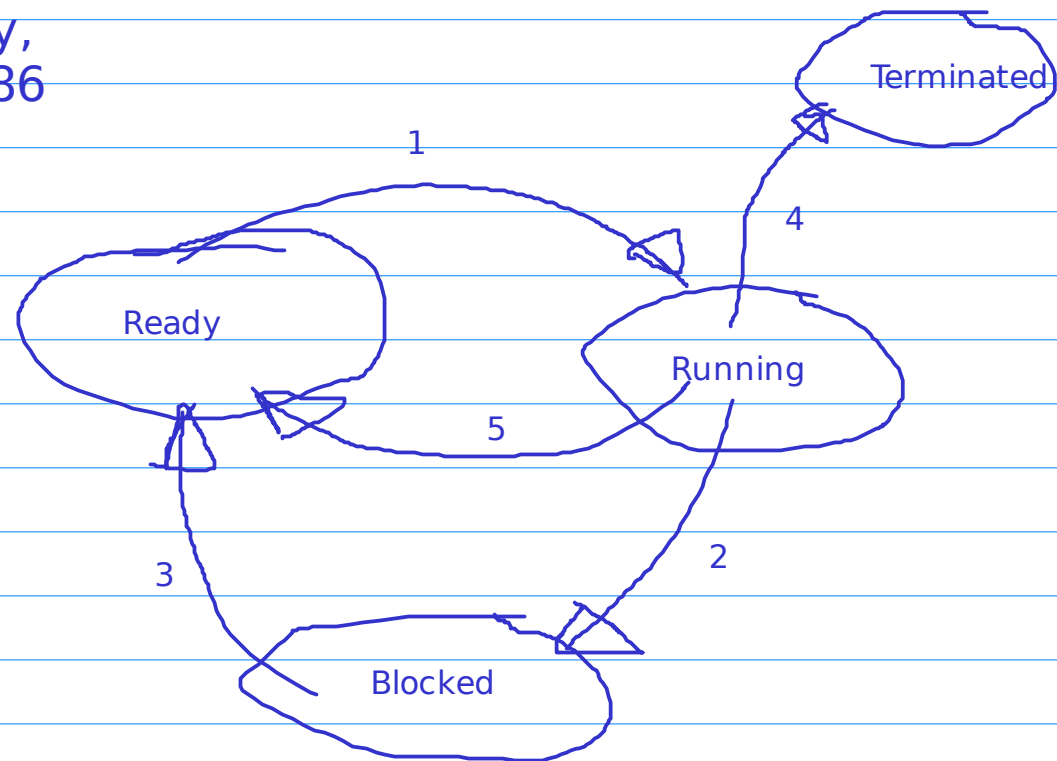
child process -- fork  
multi threading

multiplexing of CPU, under preemptive, non preemptive conditions  
context switching  
- context saving  
- context loading

context save area in memory,  
typically on top of stack in x86

Every user process  
will have independent  
stack, (independent  
address space beyond  
address space)

Threads --  
Private stack for each



Each thread/process, will have independent/private stack in userspace, as well as private stack for each thread/process in kernel space

kernel code execution can be blocked/preempted, due to private stacks -- process context of kernel execution  
-- kernel code running in process context

sometimes kernel codes in non process context (interrupt context)  
-- not associated with any user process  
-- no private/independent stack  
-- blocking/preemption not allowed

concurrency -- physical/true parallelism  
-- logical/pseudo parallelism

Task driven parallelism -- each thread does diff task  
Data driven parallelism -- same task, but operating on diff data

Userspace threading:-

pthread_create	pthread_self
pthread_join	pthread_cancel
	pthread_equal
	pthread_yield

struct task\_struct ==> attributes of each process  
current macro ==> address of task\_struct var for active process  
e.g. current->pid, represents pid of active process

init\_task, task\_struct variable for init process  
next\_task, task\_struct variable for next process in list  
list\_for\_each

-----  
Race conditions

Critical Section

Mutual Exclusion -- access to shared resource by only  
[competition] once process/thread at a time  
only can execute critical section

Solutions:-

- \* Semaphores
- \* Mutex
- \* Spinlocks

Synchronization

[cooperation, dependency, sequencing, prod/cons]

- \* Semaphores
- \* Condition vars (uspace) Wait Conditions (kspace)

val=10

val++	val--
r0 <-- val	r1 <-- val
r0 = r0 + 1	r1 = r1 - 1
r10 --> val	r1 --> val

11                      9

-----

sem.val=1

lock	lock
//critical	//critical
unlock	unlock

-----

sem.val=0

//prod:add      lock(sem)

unlock(sem)      //cons:remove

Semaphore:-

- kernel level data structure
- integer value, process/thread Q

lock/down/acquire:-

- A - if val>0, val--, go ahead
- B - if val==0, block current  
add to waiting Q

unlock/up/release:-

- C - if Q is not empty, allow any  
one waiting process/thread  
to resume
- D - if Q is empty, val++

-----

Mutex vs Semaphore


Salient features/characteristics  
of Mutex

- \* ownership applicable
- \* unlocking before locking  
unlocking more than once  
not allowed
- \* two state - T/F, unlocked/locked
- \* strictly mutual exclusion

semaphore/mutex: lock      -- block process  
                              unlock      -- resume/unblock process  
                              involves context switching

Spinlock -- busyloop technique, to avoid context switching

flag=0

P1		P2
lock : while(flag);		while(flag);
flag=1;		flag=1;
 //critical		 //critical
 flag=0;		 flag=0;

Spinlocks/busyloops are meaningful for multicore (SMP) only

H/w supported atomic instructions, (disable interrupts + bus locking)  
e.g XCHG in x86, SWP in ARM

	flag=0, reg=1	
lock:- while(XCHG(reg,flag));		lock:- while(XCHG(reg,flag));
 //critical		
unlock:- flag=0		flag=0

SMP:- Semaphore/Mutex Spinlock

compare critical section length (vs) context switching time

-----  
static DEFINE\_SEMAPHORE(s1);

down\_interruptible(&s1);

//for loop

up(&s1);

-----  
static DEFINE\_MUTEX(m1);

mutex\_lock(&m1);

//for loop

mutex\_unlock(&m1);

-----  
static DEFINE\_SPINLOCK(s1);

spin\_lock(&s1);

val++; //val--;

spin\_unlock(&s1);

## Synchronization with Semaphores:-

```
struct semaphore s2; //static DEFINE_SEMAPHORE(s2);
```

init:-

```
sema_init(&s2,0);
```

Thread-A, before for loop:-

```
down_interruptible(&s2);
```

Thread-B, end of for loop:-

```
up(&s2);
```

-----

```
wait_queue_head_t w1;  
int buflen=0;
```

```
init:-  
    init_waitqueue_head(&w1);
```

```
Thread-A:-  
    wait_event_interruptible(w1, (buflen > 0) );
```

```
Thread-B:-  
    buflen++;  
    wake_up_interruptible(&w1);
```